

Planning Under Uncertainty

Modelling and Solving POMDP for Robot Navigation

Edwin, Niranjan, Adhithiyan

RBE — WPI

27th November 2023

Introduction

This presentation shall cover the following


- ① Background and Motivation
- ② Problem Statement
- ③ Plan of Action
- ④ Tools
- ⑤ Evaluation metrics
- ⑥ Task Delegation
- ⑦ Current Environment and Proposed Framework
- ⑧ Observation and Preliminary Results
- ⑨ Future Direction

Background - Why MDPs?

- Classical path planning - an example [1]
 - Plan a sequence from start to goal - flawless execution
 - Introduce uncertainty in action - collide in narrow path




- Incorporating this uncertainty in our model - MDP
 - Markov Decision Processes - states, actions and rewards
 - Learn a strategy or policy - put agent in good value states for success
 - Example [1] - not greedy approach - risk in colliding
 - Robot rewarded for reaching goal - penalized for colliding
 - Takes a detour - more successful
 - But considers states - completely observable




Background - Why POMDPs?

- In real world - observations of states noisy - due to imperfect sensors, etc.
- Need to incorporate uncertainty in states - POMDP
 - Partially Observable Markov Decision Processes
 - Belief states instead of states
 - Probability distribution - states robot believes to be in
 - New dimension to decision making - information gathering
 - Example [1] - robot unaware of its orientation but knows initial location
 - Good strategy - first move to a corner
 - Depending on the corner - moves accordingly




Background - Problem Statement

- Robot model - Dubin's car - to reach goal configuration in a cluttered environment
- An application - steering medical needle to required tissue [2]
 - Modeled as an MDP
 - More success by avoiding a narrow, riskier path




- Objective - capture the uncertainty in states also - POMDP
- Uncertainty - robot location, environment map [3]
- Solve the POMDP - a good enough policy to reach desired configurations with more success

Background - POMDP Solvers




- Optimal policy - outputs best action to take for each possible belief state
 - After computing for specific horizon - piecewise continuous value function
 - Each region of belief space - specific action optimal
- POMDP as a problem is exponential in nature, deeper you go.
 - Curse of History
 - Curse of Dimensionality
- We will be evaluating some tree-based POMDP solvers to find good policy for the problem at hand

Plan of Action - DESPOT




- Some solvers considered for our work
 - DESPOT [4]
 - POMCP
- DESPOT works by randomly initializing a set of scenarios and then only exploring the belief space that consists of those scenarios.
- This reduces the complexity from $O(|A|_D |Z|_D)$ to $O(|A|_D K)$ nodes, where $|A|$ and $|Z|$ are the sizes of the action set and the observation set

Current Environment



- The current environment - 10×10 grid world
- Includes various scenarios of start and stop positions
 - Each state encoded as SSOO
 - SS \rightarrow Coordinate ID (79 \rightarrow 8th row 10th column)
 - OO \rightarrow Orientation ID (02 \rightarrow 90 degrees)

Dubin's Car in Grid World



- Reaching goal configuration from start using only few turns of constant curvature (LSL, RSR, RSL, etc.) under certain conditions
- Cannot travel in reverse direction
- In the grid world
 - Turns of only $\frac{\pi}{4}$ at a time allowed
 - Steps of only one grid jump allowed
 - Cannot turn without moving forward
 - Example car facing north
 - Can only go north, northwest, or northeast by one step

Proposed Framework

- **States**

- Each cell in the 10×10 grid - valid states

- **Actions**

- Normal grid traversal - 5 actions: up, down, left, right, halt
- Dubin's car - 4 actions: go straight, steer left, steer right, halt

- **Observations**

- Agent-Environment state as perceived by the agent - noisy

- **Rewards**

- +40 to state-action pairs that result in goal state
- -1 for every other valid state-action
- -10 if agent tried to leave grid

Tools and repositories used for this project are listed below:


- Python_Robotics
 - Used for visualizing and planning for Dubins Car
- Pypomdp
 - Used for solving POMDP using POMCP
 - Code for visualizing and plotting the belief state
- Despot API
 - Used to solve POMDP in pomdpx format
- APPL toolkit(SARSOP)
 - Used pomdp_convert function to convert /.POMDP to /.pomdpx format

Evaluation Metrics

- Although POMDP solvers powerful for planning
 - Notorious for high computational complexity
- Plan
 - Use tree-based solvers to get approximate optimal policies
 - Comparison between POMCP, DESPOT solvers
- Parameters
 - Online Time
 - Time taken to reach the goal and end runtime
 - Necessary to make sure the solution is not intractable.
 - Task Completion
 - Consider the probability of the robot reaching the goal
 - Evaluate in multiple scenarios


Methodology

- Problem Formulation
 - Problem was written in the /.POMDP format
 - State Transition probabilities, Observation Probabilities were defined
 - File was converted to /.POMDPX format to use different solvers
- Defining Observation Probabilities
 - Added noise distribution around neighbouring cells
 - State uncertainty decreases as we go down
- Strategy
 - Derived idea from Light-Dark method
 - Expectation - Agent moves to state with less uncertainty and gathers information




Observation - POMCP

- Agent starts at state 2 and aims to reach state 65.



((a)) Belief States




((b)) Visualization

Figure: Results from POMCP


Trajectory

- '3', '13', '23', '23', '23', '22', '22', '32', '32', '32', '32', '33', '32', '32', '33', '43', '53', '63', '64', '65', '65', '65', '65', '65', '65', '65'


Trajectory Visualization - POMCP (sort of)




((a)) Initial Belief State



((b)) After a while



((c)) After a long while




((d)) Finally...


Implementation

- Problem Definition
 - Program developed to generate POMDP file
 - Both Action Uncertainty and Observation Uncertainty are encoded
- Conversion
 - The POMDP file was converted to pomdpx utilising the SARSOP pomdpx_convert file.
- Visualization
 - Extracted States and Belief from DESPOT solver

Path in Dubins



((e)) DESPOT Solver




((f)) POMCP Solver


Figure: Results of start point [6,9]

- DESPOT Trajectory: s6902, s5902, s4803, s3802, s2902
- POMCP Trajectory: '6902', '5803', '4802', '4802', '3901', '2902'

Path in Dubins



((a)) DESPOT Solver



((b)) POMCP Solver

Figure: Results of start point [9,0]

- DESPOT Trajectory: s9000, s8101, s7201, s7300, s6401, s5402, s4402, s3501, s3600, s3700, s3800, s2901
- POMCP Trajectory: '9000', '8101', '8101', '8101', '7201', '6202', '6202', '5301', '5400', '5500', '5600', '4701', '3801', '2901'

Results - Conclusion

- 1st Instance - Start(6,9)
 - DESPOT - 5 steps
 - POMCP - 6 Steps
- 2nd Instance - Start(9,0)
 - DESPOT - 12 steps
 - POMCP - 14 Steps
- POMCP tends to halt at various positions - increased number of steps and computation time.
- Time taken by POMCP(mins) to complete the task was exceedingly large than DESPOT(secs) Despot takes ≤ 1 min whereas POMCP takes ≥ 15 mins

Task Completion - Evaluation

- Translate from grid traversal to Dubins Car (✓)
- Visualize implemented model in DESPOT Solver instead of POMCP - trajectories (✓)
- Define Gaussian Noise distribution for more exploration (✓)
- Take Action uncertainty into account (✓)
- Extract Belief states from DESPOT after each step (✓)
- Figure out why Belief State is not making sense.

References



S. Thrun, W. Burgard, and D. Fox, *Probabilistic Robotics (Intelligent Robotics and Autonomous Agents)*. The MIT Press, 2005.



R. Alterovitz, T. Siméon, and K. Goldberg, "The Stochastic Motion Roadmap: A Sampling Framework for Planning with Markov Motion Uncertainty," in *Robotics: Science and Systems III*. The MIT Press, 01 2008. [Online]. Available: <https://doi.org/10.7551/mitpress/7830.003.0031>



M. Lauri, D. Hsu, and J. Pajarinen, "Partially observable markov decision processes in robotics: A survey," *IEEE Transactions on Robotics*, vol. 39, no. 1, pp. 21–40, 2023.



N. Ye, A. Somani, D. Hsu, and W. S. Lee, "Despot: Online pomdp planning with regularization," *The Journal of artificial intelligence research*, vol. 58, pp. 231–266, 2017.